

Explaining AI in Finance: Past, Present, Prospects

NO AUTHOR

The proliferation of Artificial Intelligence (AI) in finance has revolutionized operational strategies. Yet, the challenge of interpretability, encapsulated in the concept of explainable AI (XAI), remains pivotal for its comprehensive adoption. This paper navigates the evolution of AI in finance, demonstrating the role and potential of XAI within this transformation. We trace the journey from primitive statistical methods to modern machine learning, addressing traditional models' inability to interpret intricate non-linear relationships in financial variables. We delve into XAI's role in prevalent financial applications like credit scoring and fraud detection, demonstrating how methods like Shapley values provide superior, nuanced inference compared to traditional linear regression models in complex scenarios. A vital discussion is the implications of upcoming EU regulations, emphasizing the need for further XAI research to bridge the gap between intricate model decision-making and human interpretability. This narrative is bolstered by several simulated examples, comparing the explanatory power of Shapley values with conventional regression coefficients in the context of asset pricing, demonstrating the effectiveness of XAI in finance. In conclusion, the paper accentuates the crucial role of XAI in enhancing trust in AI systems within finance, promoting more informed and responsible decision-making. Through this, we underline the promise of XAI as a facilitator of transparency, reliability, and ultimately, the wider acceptance of AI in the finance industry.

Introduction

The rapid digital transformation of the finance industry over the past few decades has been predominantly driven by the integration of Artificial Intelligence (AI) and machine learning technologies. These technologies have heralded a new era in finance, catalysing innovations in trading, risk management, fraud detection, customer service, and many other areas, bringing significant changes to business models, operations, and services (Arner, Barberis, and Buckley

2020). Today, financial institutions leverage these advanced technologies to generate insights, automate processes, and improve decision-making.

Despite the revolutionary potential of AI, its application in finance is not devoid of challenges. For example, a significant issue arises around the transparency and interpretability of AI decision-making, often described as the ‘black box’ problem. This term refers to the difficulty in understanding how complex AI and machine learning models arrive at their decisions (Dhar 2018). This opacity presents substantial ethical, legal, and practical challenges, especially in an industry as regulated and risk-averse as finance (Bhatt, Xiang, and Sharma 2020).

In response to these concerns, Explainable Artificial Intelligence (XAI) has emerged to make AI’s decision-making process more transparent and understandable to human users (Molnar 2020). This development is particularly crucial in the financial sector, where understanding the rationale behind decisions can have enormous implications for trust, compliance, risk management, and customer satisfaction.

This paper aims to critically analyse the role of XAI in finance, tracing its historical development, examining its current applications, and exploring its prospects. By providing a comprehensive review of XAI in the context of finance, we hope to shed light on its importance and potential for the industry while addressing ongoing challenges and areas for future research.

Literature review

Past: Early Application of AI in Finance

The integration of AI in finance has a history that dates back to the latter half of the 20th century. Initially, financial institutions deployed rule-based systems for various functions, such as automated trading and risk analysis. These were the earliest forms of AI applied to finance and were quite simplistic compared to today’s advanced systems.

These early AI systems were based on sets of pre-programmed rules, often developed by human experts, and were typically deterministic. Given the same input, these systems would always provide the same output. They were transparent and easily interpretable because they followed clearly defined, pre-set rules (Gomber et al. 2018).

However, the effectiveness of rule-based systems was limited by their rigid, inflexible design. These systems were not designed to learn from new data or adapt to changing conditions, making them less useful in the dynamic world of finance, which is characterised by evolving markets, changing regulatory landscapes, and unpredictable economic conditions (Dhar 2018).

The desire for more adaptive, responsive systems led to the developing of machine learning algorithms. Machine learning represented a significant advancement in the field of AI. Unlike

rule-based systems, machine learning models could learn from data, identify patterns, and make predictions or decisions based on those patterns. Machine learning algorithms found applications in various areas of finance, from predicting stock prices to identifying fraudulent transactions (Athey 2021).

The transition from rule-based systems to machine learning models marked a pivotal shift in AI's role in finance. However, this transition was not without its challenges. Machine learning models, particularly more complex ones like neural networks, introduced a level of opacity and complexity that made it difficult for human users to understand how they made decisions (Rudin 2019). This lack of transparency and interpretability in machine learning systems became known as the 'black box' problem and formed the backdrop against which the field of explainable AI emerged.

In summary, the past of AI in finance was marked by the transition from transparent but inflexible rule-based systems to powerful but opaque machine learning models. The need to address the transparency issues introduced by machine learning has led to the development of explainable AI, the current and future state of which will be discussed in the following sections.

Present: The Rise of Machine Learning and XAI

In the present landscape, machine learning models, such as neural networks, decision trees, and support vector machines, have become integral parts of financial institutions. These models perform various tasks, including credit scoring, fraud detection, algorithmic trading, portfolio optimisation, and customer segmentation (Athey 2021).

Despite their efficiency and sophistication, these models often work as 'black boxes,' where the internal decision-making process is obscured from the users. This opaqueness can pose considerable challenges. For example, it hinders human users, such as loan officers or portfolio managers, from understanding and trusting the model's decisions. In addition, it poses problems for accountability and compliance on a regulatory level, especially in jurisdictions where decisions affecting individuals must be explainable (Rudin 2019).

Enter the Explainable Artificial Intelligence (XAI) field, which seeks to make AI's decision-making process more transparent and interpretable. Various techniques fall under the umbrella of XAI, and these can be broadly classified into model-specific methods and model-agnostic methods (Molnar 2020).

Model-specific methods, such as coefficient interpretation in linear regression or feature importance in decision trees, provide insights into how these specific models operate. However, their application is limited to the particular model types for which they were developed (Molnar 2020).

Model-agnostic methods, on the other hand, can be applied to any machine-learning model. They seek to provide explanations for individual predictions regardless of the complexity or

type of the underlying model. Examples of these techniques include Local Interpretable Model-Agnostic Explanations (LIME) and SHapley Additive exPlanations (SHAP). LIME offers explanations by approximating the prediction of a complex model with a simpler, locally-fitted model around the prediction point (Ribeiro, Singh, & Guestrin, 2016). Meanwhile, SHAP allocates each feature’s contribution to the prediction for individual data points based on concepts from cooperative game theory (Lundberg and Lee 2017).

Yet, despite these advancements, achieving true explainability in AI remains a significant challenge. Many of these methods provide post-hoc explanations, which attempt to interpret the model’s behaviour after training. This process often involves a trade-off between accuracy and interpretability, with more complex models offering greater accuracy but less interpretability (Bhatt, Xiang, and Sharma 2020).

Moreover, explainability is not just a technical problem but also a human-centred one. The effectiveness of an explanation largely depends on the recipient’s perspective and the context in which it is given (Miller, 2019). For example, a satisfactory explanation to a data scientist might be incomprehensible to a loan officer or a customer, indicating that the development of XAI needs to consider the human factors of understandability and trust.

In conclusion, the present state of XAI in finance is marked by considerable advancements and ongoing challenges. The shift towards more transparent and interpretable AI models is underway, with various methods being developed and applied. However, balancing models’ complexity (and, thereby, performance) and their interpretability remains a significant hurdle. As we look towards the future, it is crucial that these challenges are addressed and XAI continues to evolve to meet the demands of transparency and interpretability in the financial industry.

Prospects: The Future of XAI in Finance

As we move into the future, the demand for transparency and interpretability in AI systems within the finance sector is expected to grow. The future advancements and challenges in the field of XAI reflect this.

One significant direction for future research and development is integrating explainability directly into the model-building process rather than treating it as an afterthought. This approach, often called intrinsic explainability, involves building naturally interpretable models, such as explainable boosting machines or interpretable decision sets (Lakkaraju, Bach, and Leskovec 2016; Lou, Caruana, and Gehrke 2012). Developing such models can help mitigate the trade-off between accuracy and interpretability that characterises post-hoc explanation methods.

Furthermore, as the field of XAI evolves, it will be essential to focus on the users’ perspective. What constitutes a ‘good’ explanation can vary based on the recipient and the context. Therefore, future XAI methods should tailor answers to different users’ needs and capabilities.

They should also address how to communicate these explanations best to ensure they are understandable and valuable (Miller, 2019).

Moreover, as AI and XAI become more commonplace in finance, regulations will likely evolve to address their new challenges. For example, the European Union’s General Data Protection Regulation (GDPR) has already introduced a ‘right to explanation’, where individuals can ask for explanations of decisions made by automated systems that affect them (Goodman and Flaxman 2017). In the future, we might see more regulations requiring financial institutions to provide clear, understandable explanations for AI-based decisions.

However, implementing such regulations comes with its challenges. For example, regulators will need to define what constitutes an ‘explanation’ and a ‘decision’ in the context of AI. They will also need to set standards for how detailed and understandable these explanations must be (Edwards and Veale 2017).

In summary, the future of XAI in finance is ripe with opportunities for making AI decision-making more transparent and accountable. However, it also presents challenges that must be addressed through continued research, development, and thoughtful regulation.

Confronting Econometrics with XAI

Explainable AI (XAI) could undoubtedly play a significant role in improving the field of econometrics, which is the application of statistical methods to economic data to give empirical content to economic relations.

Traditionally, econometric models have been designed to be inherently interpretable, as they often depend on linear relationships and other simplistic assumptions to ensure that the model’s parameters can be easily interpreted. However, these assumptions can be limiting, as they might not fully capture the complexities of real-world economic phenomena.

With the advent of machine learning, econometricians have created models that can learn complex patterns from data, leading to more accurate predictions. But the drawback is that these models are often ‘black boxes,’ making it difficult to understand how they make decisions.

This is where XAI comes in. By using techniques developed under the umbrella of XAI, econometricians could make their machine-learning models more transparent, allowing them to understand how each input variable contributes to the model’s predictions. This increased transparency could make these models more acceptable to economists and policymakers, who need to understand the decision-making process to make informed decisions.

For instance, SHapley Additive exPlanations (SHAP) and Local Interpretable Model-agnostic Explanations (LIME) could be applied to make machine learning models more interpretable. These techniques explain individual predictions, helping to understand the contributions of different features to the model’s output.

In addition to improving model transparency, XAI could help address some statistical challenges in econometrics. For example, XAI could help econometricians understand the variable importance in their models, which could help address issues related to multicollinearity, where independent variables in a regression model are highly correlated.

In conclusion, XAI holds considerable potential in improving econometrics, mainly as the field increasingly incorporates machine learning models. Furthermore, by making these models more transparent, XAI can help econometricians and policymakers better understand and trust the predictions derived from these models, leading to better decision-making.

Shapley values versus OLS regression coefficients

In a linear regression model, let's explore the analogy between Shapley values and the coefficients.

In a linear regression model, each predictor variable is assigned a coefficient representing its partial effect on the outcome variable, controlling for all other predictors. The coefficient can be interpreted as the expected change in the outcome variable for a one-unit change in the predictor, holding all other predictors constant.

Analogously, the Shapley value for a player in a cooperative game represents the average contribution of the player to the worth of all possible coalitions that the player can be a part of. The Shapley value considers all possible ways the coalition can be formed and averages over them.

In both cases, the goal is to fairly distribute some total quantity (the total worth of the grand coalition in a cooperative game or the total variance of the outcome variable in a linear regression model) among different contributors (the players in a cooperative game, or the predictor variables in a linear regression model).

However, there are significant differences as well. While linear regression assumes a specific linear and additive form for the relationship between predictors and the outcome, Shapley values make no such assumption. Shapley values can handle any type of game, including non-cooperative games and games with complex interactions between players.

Also, the computation of Shapley values considers all possible orders in which players can join the coalition, reflecting the idea that a player's contribution may depend on which other players are already in the coalition. In contrast, linear regression coefficients are typically computed using a method (like ordinary least squares) that does not consider different orders of entering the predictors into the model.

In the context of explainable AI, the Shapley value concept has been applied to machine learning models to compute the contribution of each feature to the prediction for a particular instance. This can provide more nuanced and reliable interpretations than simply looking at

the coefficients of a linear model, especially for complex models that capture non-linear and interactive effects.

Coalition game and Shapley values

Let's consider a simple cooperative game involving three players: A, B, and C, to fix ideas. A characteristic function v gives the worth of each coalition of players:

- $v(\{\}) = 0$ (worth of the empty coalition)
- $v(\{A\}) = 100$
- $v(\{B\}) = 200$
- $v(\{C\}) = 300$
- $v(\{A, B\}) = 400$
- $v(\{A, C\}) = 500$
- $v(\{B, C\}) = 600$
- $v(\{A, B, C\}) = 800$

We want to distribute the total worth of the grand coalition ($v(\{A, B, C\}) = 800$) among the players in a way that reflects their contribution to the coalition.

The Shapley value is one way to do this. For each player, it computes the average marginal contribution of the player to all possible coalitions. This is done by considering all permutations of the players and, for each permutation, adding up the marginal contributions of the player when they join the coalition.

Confronting regression coefficients using XAI

In a linear regression model, each predictor variable is assigned a coefficient representing its partial effect on the outcome variable, controlling for all other predictors. The coefficient can be interpreted as the expected change in the outcome variable for a one-unit change in the predictor, holding all other predictors constant.

Analogously, the Shapley value for a player in a cooperative game represents the average contribution of the player to the worth of all possible coalitions that the player can be a part of. The Shapley value considers all possible ways the coalition can be formed and averages over them.

In both cases, the goal is to fairly distribute some total quantity (the total worth of the grand coalition in a cooperative game or the total variance of the outcome variable in a linear regression model) among different contributors (the players in a cooperative game, or the predictor variables in a linear regression model).

However, there are significant differences as well. While linear regression assumes a specific linear and additive form for the relationship between predictors and the outcome, Shapley

values make no such assumption. Shapley values can handle any type of game, including non-cooperative games and games with complex interactions between players.

Also, the computation of Shapley values considers all possible orders in which players can join the coalition, reflecting the idea that a player's contribution may depend on which other players are already in the coalition. In contrast, linear regression coefficients are typically computed using a method (like ordinary least squares) that does not consider different orders of entering the predictors into the model.

In the context of explainable AI, the Shapley value concept has been applied to machine learning models to compute the contribution of each feature to the prediction for a particular instance. This can provide more nuanced and reliable interpretations than simply looking at the coefficients of a linear model, especially for complex models that capture non-linear and interactive effects.

A finance example

To Steelman, the case here is a simulated example. We'll use the `sklearn` and `shap` libraries in Python to fit a linear regression model and a more complex model (Random Forest) to the simulated data. Then we'll use the `shap` library to compute SHAP values for the Random Forest model.

Let's consider a simple three-factor model (size, value, and momentum factors) to predict asset returns.

First, let's install the necessary packages (if not already installed):

Now let's import the necessary packages and generate some simulated data: g

```
```{python}
import numpy as np
import pandas as pd
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
import warnings
warnings.filterwarnings("ignore")
import shap

Set a seed for reproducibility
np.random.seed(0)

Generate simulated factor values
n = 1000 # number of assets
size = np.random.normal(0, 1, n)
```



```

value = np.random.normal(0, 1, n)
momentum = np.random.normal(0, 1, n)

Generate simulated asset returns
The true model has coefficients [0.1, 0.2, 0.3] for the factors
We also add some noise
returns = 0.1*size + 0.2*value + 0.3*momentum + np.random.normal(0, 0.1, n)

Put data in a DataFrame
data = pd.DataFrame({
 'Size': size,
 'Value': value,
 'Momentum': momentum,
 'Return': returns
})

```

Next, let's fit a linear regression model to the data:

```

{python}
X = data[['Size', 'Value', 'Momentum']]
y = data['Return']

lin_reg = LinearRegression().fit(X, y)

Print coefficients
print('Linear Regression Coefficients:')
print('Size:', lin_reg.coef_[0])
print('Value:', lin_reg.coef_[1])
print('Momentum:', lin_reg.coef_[2])

```

```

Linear Regression Coefficients:
Size: 0.10091701879836647
Value: 0.20185690366007059
Momentum: 0.30119085670658874

```

Now let's fit a Random Forest model to the data and compute SHAP values:

```

{python}
Fit model

```

```

rf = RandomForestRegressor(n_estimators=100, random_state=0).fit(X, y)

Explain model predictions using SHAP
explainer = shap.Explainer(rf, X)
shap_values = explainer(X)

Print mean absolute SHAP values for each feature
print('\nMean Absolute SHAP Values:')
print('Size:', np.mean(np.abs(shap_values.values[:, 0])))
print('Value:', np.mean(np.abs(shap_values.values[:, 1])))
print('Momentum:', np.mean(np.abs(shap_values.values[:, 2])))
```

```

```

85%|===== | 847/1000 [00:11<00:01]
93%|===== | 929/1000 [00:12<00:00]

```

```

Mean Absolute SHAP Values:
Size: 0.0662061209562134
Value: 0.15558826181291682
Momentum: 0.23316567312446296

```

In this simulated example, the actual model is linear, so the linear regression coefficients should be close to the actual values, and the mean absolute SHAP values for the Random Forest model should also reflect the relative importance of the factors.

However, in real-world scenarios where the relationship between factors and asset returns might be non-linear or involve interactions, the Random Forest model could provide better predictive accuracy, and the SHAP values could offer a more nuanced understanding of feature contributions.

Let's modify the above example to include non-linear and interaction effects. In this modified example, the true polynomial model includes square terms and an interaction term.

```

```{python}
Generate simulated asset returns with non-linear and interaction effects
returns = 0.1*size + 0.2*value + 0.3*momentum + 0.4*size**2 + 0.5*value**2 + 0.6*momentum*

Put data in a DataFrame
data = pd.DataFrame({
 'Size': size,
 'Value': value,

```

```

 'Momentum': momentum,
 'Return': returns
 })

X = data[['Size', 'Value', 'Momentum']]
y = data['Return']
```

```

Next, let's fit a linear regression model to the data:

```

```{python}
lin_reg = LinearRegression().fit(X, y)

Print coefficients
print('Linear Regression Coefficients:')
print('Size:', lin_reg.coef_[0])
print('Value:', lin_reg.coef_[1])
print('Momentum:', lin_reg.coef_[2])
```

```

```

Linear Regression Coefficients:
Size: 0.10773946297944371
Value: 0.2764299938154271
Momentum: 0.16456043298921597

```

Now let's fit a Random Forest model to the data and compute SHAP values:

```

```{python}
Fit model
rf = RandomForestRegressor(n_estimators=100, random_state=0).fit(X, y)

Explain model predictions using SHAP
explainer = shap.Explainer(rf, X)
shap_values = explainer(X)

Print mean absolute SHAP values for each feature
print('\nMean Absolute SHAP Values:')
print('Size:', np.mean(np.abs(shap_values.values[:, 0])))
print('Value:', np.mean(np.abs(shap_values.values[:, 1])))
print('Momentum:', np.mean(np.abs(shap_values.values[:, 2])))
```

```

```
83%|=====| 830/1000 [00:11<00:02]
91%|=====| 907/1000 [00:12<00:01]
99%|=====| 986/1000 [00:13<00:00]
```

Mean Absolute SHAP Values:
Size: 0.3687139875766077
Value: 0.458121562987211
Momentum: 0.4565856202986719

In this case, the linear regression coefficients fail to capture the non-linear and interaction effects. The Random Forest model, which can capture these complex relationships, should provide better predictive accuracy. The SHAP values for the Random Forest model offer a more nuanced understanding of feature contributions, taking into account the non-linear and interaction effects.

Let's create a scenario where the relationship between the factors and the return is non-linear. This is where Shapley values can potentially provide a more nuanced view than linear regression coefficients.

We will use a similar setup but introduce some interaction and non-linearity in how the factors influence the return.

```
```{python}
import numpy as np
import pandas as pd
from sklearn.ensemble import RandomForestRegressor
import shap

np.random.seed(0)

Create a sample dataframe with 2 factors and a response
n = 1000
Factor1 = np.random.rand(n)
Factor2 = np.random.rand(n)
Noise = 0.1*np.random.randn(n)

Assume factors interact and have a non-linear relationship with the response
Return = 3*Factor1**2 + 5*np.exp(Factor2) + 2*Factor1*Factor2 + Noise

data = pd.DataFrame({
 'Return': Return,
 'Factor1': Factor1,
```

```

 'Factor2': Factor2
})

Fit a random forest regression model
model = RandomForestRegressor(n_estimators=100)
model.fit(data[['Factor1', 'Factor2']], data['Return'])

Calculate Shapley values
explainer = shap.Explainer(model, data[['Factor1', 'Factor2']])
shap_values = explainer(data[['Factor1', 'Factor2']])

Print the mean absolute Shapley values for each feature
print("\nMean absolute Shapley Values:")
print("Factor1: ", np.mean(np.abs(shap_values.values[:,0])))
print("Factor2: ", np.mean(np.abs(shap_values.values[:,1])))
```

```

```
RandomForestRegressor()
```

```
Mean absolute Shapley Values:
```

```
Factor1: 1.0587623348271473
```

```
Factor2: 2.4884202639370416
```

In this example, we use a random forest regressor to handle non-linear and interaction effects among factors. However, this model is more complex than a simple linear regression, and the “coefficients” or feature importances do not clearly measure each factor’s contribution, especially when there is interaction and non-linearity.

After training the model, we compute Shapley values for each observation. Shapley values have the potential to provide a more nuanced view of the importance of each factor, as they account for the effect of each feature in all possible subsets of features. This can be especially helpful in scenarios where factors have interaction effects or contribute to the response in a non-linear way.

As in the previous example, we calculate each factor’s mean absolute Shapley value across all observations. In a situation with interaction and non-linearity, these values provide a clearer picture of the average contribution of each factor to the return across all possible orders of inputting the factors into the model.

The ability of Shapley values to handle interaction and non-linear effects is a powerful feature. It can be crucial in scenarios where linear regression coefficients fail to understand feature

importance clearly. This is especially true in finance, where the relationship between variables can often be complex and non-linear.¹

Conclusion

As the application of AI in the financial sector continues to grow, so does the need for transparency and interpretability in AI decision-making. This paper has provided a critical analysis of the role of Explainable Artificial Intelligence (XAI) in finance, tracing its evolution from the past, analysing its present state, and exploring its prospects for the future.

Historically, the financial sector witnessed a transition from rule-based AI systems to complex machine learning models. However, while these new models offered improved performance, they also introduced challenges concerning transparency and interpretability. This shift set the stage for the emergence of XAI, aiming to shed light on the ‘black box’ of AI.

Various XAI techniques, both model-specific and model-agnostic, have been developed and implemented in financial institutions. These techniques have shown promise in improving the interpretability of AI systems, enhancing user trust, and facilitating regulatory compliance. However, achieving a balance between model performance and interpretability remains a challenge. Additionally, the field is grappling with questions about what constitutes a ‘good’ explanation and how to communicate these explanations to different users effectively.

Looking to the future, XAI is expected to continue to evolve, integrating explainability directly into the model-building process and tailoring explanations to users’ needs and contexts. At the same time, it is anticipated that regulations around AI and XAI will develop further, presenting new challenges and opportunities for the field.

In conclusion, XAI represents a critical frontier in integrating AI within finance. As we continue to navigate this frontier, we must remain committed to advancing the field, improving the transparency and interpretability of AI systems, and addressing the ethical, legal, and practical challenges that arise along the way. The journey to fully explainable AI in finance may be complex and fraught with challenges, but it is a journey that holds considerable promise for the future of the industry.

References

- Arner, D. W., J. Barberis, and R. P. Buckley. 2020. “The Evolution of Fintech: A New Post-Crisis Paradigm.” *Georgetown Journal of International Law* 47 (4): 1271–1319.
- Athey, S. 2021. “The Impact of Machine Learning on Economics.” In *The Economics of Artificial Intelligence: An Agenda*, 507–47. University of Chicago Press.

¹A docker image of the simulations can be provided upon request for replication purposes.

- Bhatt, U., A. Xiang, and S. Sharma. 2020. “Explainable Machine Learning in Deployment.” In *Proceedings of the 2020 Conference on Fairness, Accountability, and Transparency*, 648–57.
- Dhar, V. 2018. “Machine Learning—the Future of Finance.” In *Machine Learning Applications for Data Center Optimization*, 21–31. Cham: Springer.
- Edwards, L., and M. Veale. 2017. “Slave to the Algorithm? Why a ‘Right to an Explanation’ Is Probably Not the Remedy You Are Looking For.” *Duke L. & Tech. Rev.* 16: 18.
- Gomber, P., R. J. Kauffman, C. Parker, and B. W. Weber. 2018. “On the Fintech Revolution: Interpreting the Forces of Innovation, Disruption, and Transformation in Financial Services.” *Journal of Management Information Systems* 35 (1): 220–65.
- Goodman, B., and S. Flaxman. 2017. “European Union Regulations on Algorithmic Decision-Making and a” Right to Explanation.” *AI Magazine* 38 (3): 50–57.
- Lakkaraju, H., S. H. Bach, and J. Leskovec. 2016. “Interpretable Decision Sets: A Joint Framework for Description and Prediction.” In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 1675–84.
- Lou, Y., R. Caruana, and J. Gehrke. 2012. “Intelligible Models for Classification and Regression.” In *Proceedings of the 18th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 150–58.
- Lundberg, Scott M., and Su-In Lee. 2017. “A Unified Approach to Interpreting Model Predictions.” *Advances in Neural Information Processing Systems*, 4768–77.
- Molnar, Christoph. 2020. *Interpretable Machine Learning*. Lulu.com.
- Rudin, C. 2019. “Stop Explaining Black Box Machine Learning Models for High Stakes Decisions and Use Interpretable Models Instead.” *Nature Machine Intelligence* 1 (5): 206–15.